

## 5241Assignment1

Francis Zhang xz3279

February 7, 2024

## Problem 1

We simulate in a simple way. Use  $y = 3x + 2 + \text{noise}$ . And then see the MSEs for both training set and test set.

1. More flexibility  $\implies$  fits more datas. So the MSEs should be decreasing here.
2. In general, more flexibility means less  $\text{bias}^2$  and more variance. Increasing flexibility at first lowers total error a lot. However, as  $\text{bias}^2$  is approaching 0, but variance can be much more and more. So total error of training set is decreasing at first to a optimal value, and then increases. Which is **Bias-Variance Tradeoff**. Simulations are attached below. In my perspective, I'm not sure about the MSE here is in which part of the U-curve(**Bias-Variance Tradeoff**). We begin our simulation below for both 1. and 2.. (The rest of my "text" answer is on page 9.)

```
[63]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import make_pipeline

np.random.seed(1) # Set seed for reproducibility

n = 30 # Number of observations

# Generate training data
x = np.sort(np.random.uniform(-3, 3, n))
# Linear Function: y = 2 + 3x + noise
y = 2 + 3 * x + 2 * np.random.randn(n) #adding random noise, times 2 to make it
    ↪ more noisier

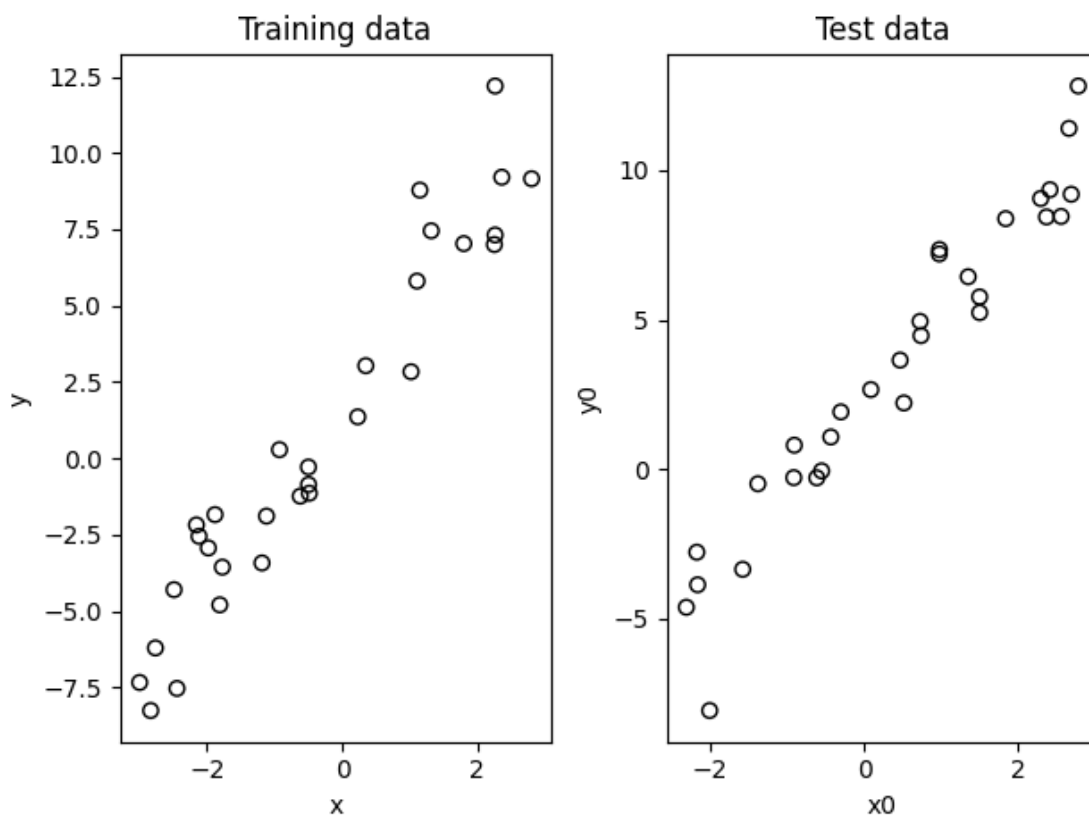
# Generate test data
x0 = np.sort(np.random.uniform(-3, 3, n))
y0 = 2 + 3 * x0 + 2 * np.random.randn(n)

# Set up a 1x2 layout for plotting
plt.subplot(1, 2, 1)
plt.scatter(x, y, marker='o', facecolors='none', edgecolors='black') # Hollow
    ↪ black circles
plt.title("Training data")
```

```
plt.xlabel("x")
plt.ylabel("y")

plt.subplot(1, 2, 2)
plt.scatter(x0, y0, marker='o', facecolors='none', edgecolors='black') # Hollow
    ↳ black circles
plt.title("Test data")
plt.xlabel("x0")
plt.ylabel("y0")

plt.tight_layout()
plt.show()
```



```
[85]: # Fit linear regression model to training data
lm_1 = LinearRegression()
lm_1.fit(x.reshape(-1, 1), y)

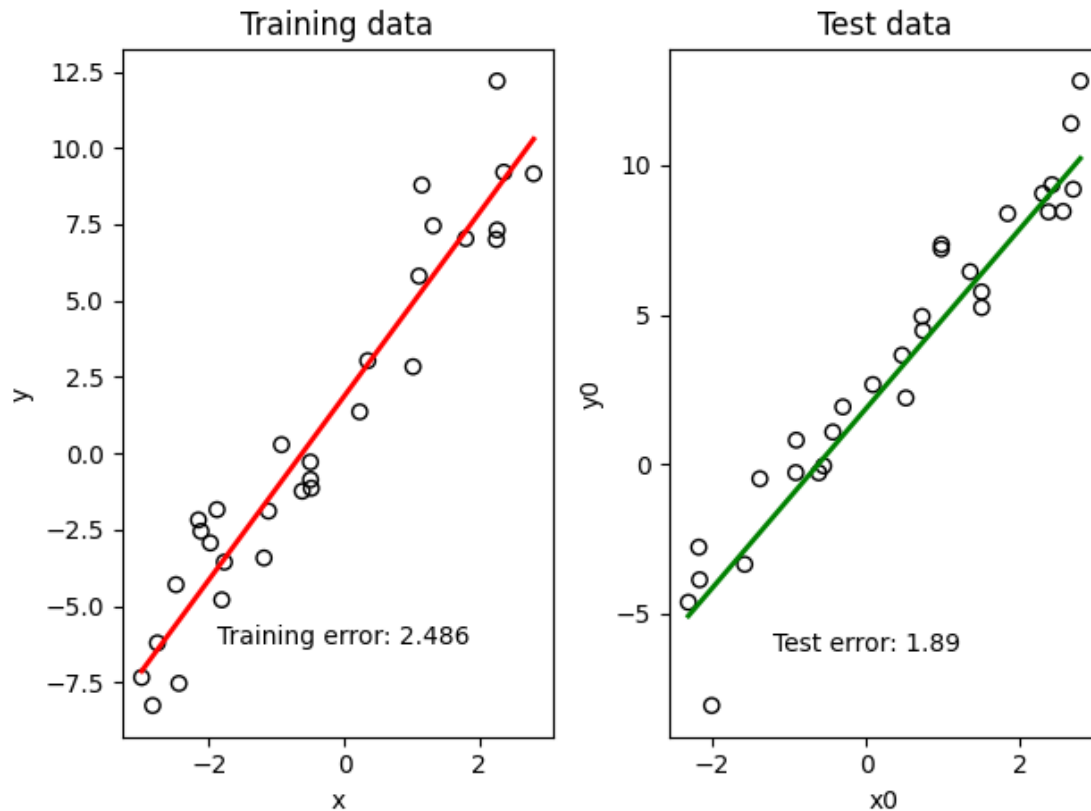
# Predictions on training data
yhat_1 = lm_1.predict(x.reshape(-1, 1))
train_err_1 = np.mean((y - yhat_1)**2)
```

```
# Predictions on test data
y0hat_1 = lm_1.predict(x0.reshape(-1, 1))
test_err_1 = np.mean((y0 - y0hat_1)**2)

# Set up a 1x2 layout for plotting
plt.subplot(1, 2, 1)
plt.scatter(x, y, marker='o', facecolors='none', edgecolors='black') # Scatter plot of training data
plt.plot(x, yhat_1, color='red', linewidth=2) # Regression line
plt.title("Training data")
plt.xlabel("x")
plt.ylabel("y")
plt.text(0, -6, f"Training error: {round(train_err_1, 3)}", ha='center', va='center', color='black')

plt.subplot(1, 2, 2)
plt.scatter(x0, y0, marker='o', facecolors='none', edgecolors='black') # Scatter plot of test data
plt.plot(x0, y0hat_1, color='green', linewidth=2) # Regression line
plt.title("Test data")
plt.xlabel("x0")
plt.ylabel("y0")
plt.text(0, -6, f"Test error: {round(test_err_1, 3)}", ha='center', va='center', color='black')

plt.tight_layout()
plt.show()
```



```
[86]: lm_1.intercept_
```

```
[86]: 1.8624318882799002
```

```
[87]: lm_1.coef_
```

```
[87]: array([2.99977939])
```

```
[88]: # Fit 2th order polynomial regression model to training data
lm_2 = make_pipeline(PolynomialFeatures(degree=2), LinearRegression())
lm_2.fit(x.reshape(-1, 1), y)

# Predictions on training data
yhat_2 = lm_2.predict(x.reshape(-1, 1))
train_err_2 = np.mean((y - yhat_2)**2)

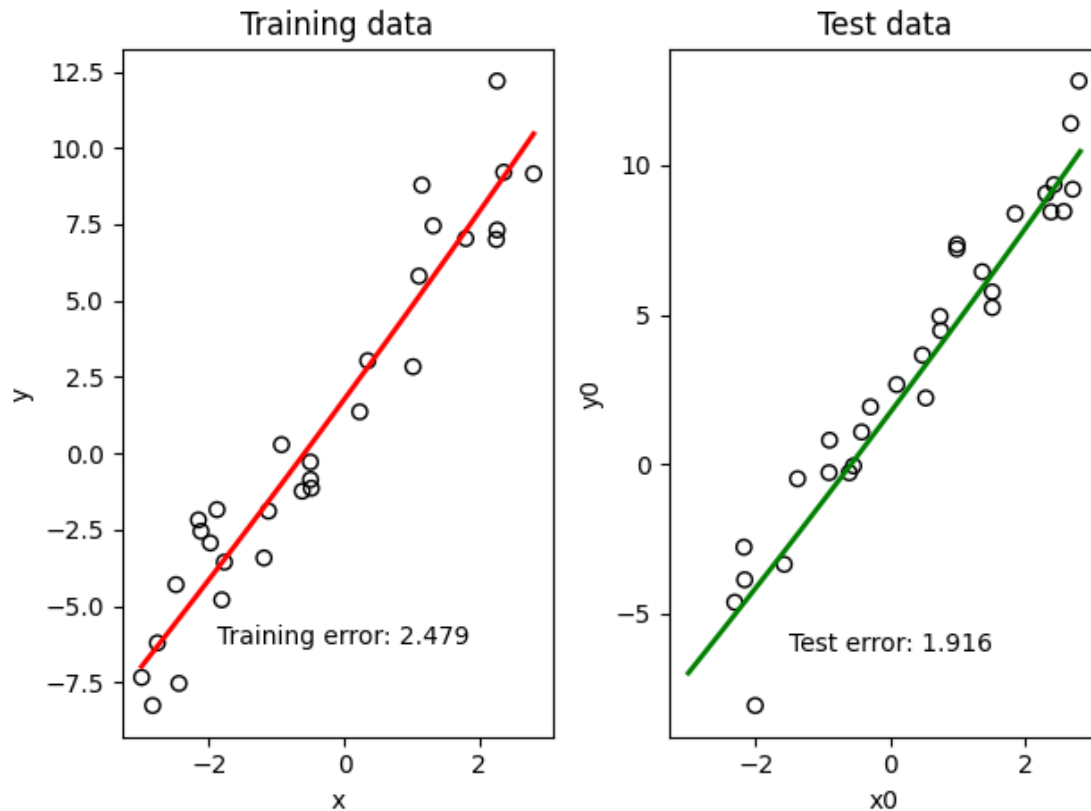
# Predictions on test data
y0hat_2 = lm_2.predict(x0.reshape(-1, 1))
test_err_2 = np.mean((y0 - y0hat_2)**2)

# Set up a 1x2 layout for plotting
```

```
plt.subplot(1, 2, 1)
plt.scatter(x, y, marker='o', facecolors='none', edgecolors='black') # Scatter plot of training data
xx = np.linspace(min(x), max(x), 100)
plt.plot(xx, lm_2.predict(xx.reshape(-1, 1)), color='red', linewidth=2) # Regression line
plt.title("Training data")
plt.xlabel("x")
plt.ylabel("y")
plt.text(0, -6, f"Training error: {round(train_err_2, 3)}", ha='center', va='center', color='black')

plt.subplot(1, 2, 2)
plt.scatter(x0, y0, marker='o', facecolors='none', edgecolors='black') # Scatter plot of test data
plt.plot(xx, lm_2.predict(xx.reshape(-1, 1)), color='green', linewidth=2) # Regression line
plt.title("Test data")
plt.xlabel("x0")
plt.ylabel("y0")
plt.text(0, -6, f"Test error: {round(test_err_2, 3)}", ha='center', va='center', color='black')

plt.tight_layout()
plt.show()
```



```
[89]: lm_2.named_steps['linearregression'].intercept_
```

```
[89]: 1.7572092032211697
```

```
[90]: lm_2.named_steps['linearregression'].coef_
```

```
[90]: array([0.          , 3.00921458, 0.03195098])
```

```
[91]: # Fit 3th order polynomial regression model to training data
lm_3 = make_pipeline(PolynomialFeatures(degree=3), LinearRegression())
lm_3.fit(x.reshape(-1, 1), y)

# Predictions on training data
yhat_3 = lm_3.predict(x.reshape(-1, 1))
train_err_3 = np.mean((y - yhat_3)**2)

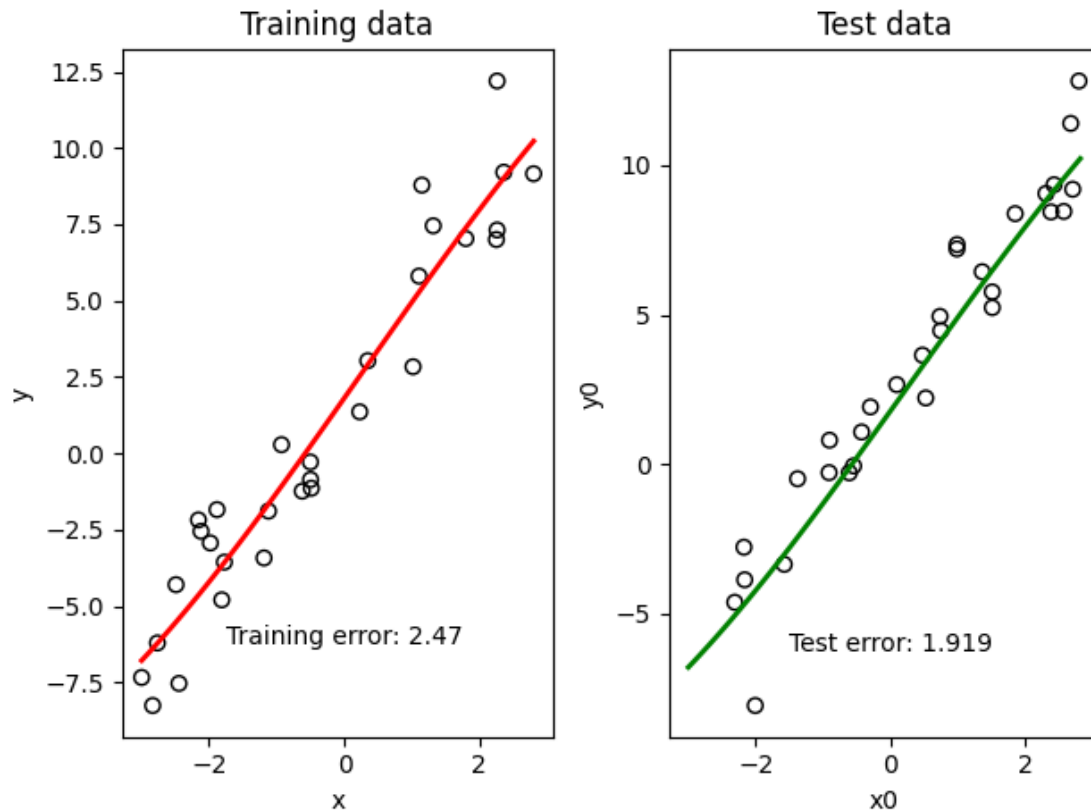
# Predictions on test data
y0hat_3 = lm_3.predict(x0.reshape(-1, 1))
test_err_3 = np.mean((y0 - y0hat_3)**2)

# Set up a 1x2 layout for plotting
```

```
plt.subplot(1, 2, 1)
plt.scatter(x, y, marker='o', facecolors='none', edgecolors='black') # Scatter plot of training data
xx = np.linspace(min(x), max(x), 100)
plt.plot(xx, lm_3.predict(xx.reshape(-1, 1)), color='red', linewidth=2) # Regression line
plt.title("Training data")
plt.xlabel("x")
plt.ylabel("y")
plt.text(0, -6, f"Training error: {round(train_err_3, 3)}", ha='center', va='center', color='black')

plt.subplot(1, 2, 2)
plt.scatter(x0, y0, marker='o', facecolors='none', edgecolors='black') # Scatter plot of test data
plt.plot(xx, lm_3.predict(xx.reshape(-1, 1)), color='green', linewidth=2) # Regression line
plt.title("Test data")
plt.xlabel("x0")
plt.ylabel("y0")
plt.text(0, -6, f"Test error: {round(test_err_3, 3)}", ha='center', va='center', color='black')

plt.tight_layout()
plt.show()
```



```
[92]: lm_3.named_steps['linearregression'].intercept_
```

```
[92]: 1.7961030966373994
```

```
[93]: lm_3.named_steps['linearregression'].coef_
```

```
[93]: array([ 0.          ,  3.14176142,  0.01900498, -0.02478674])
```

We can see that when the flexibility increases,  $MSE_{training}$  decreases and  $MSE_{test}$  increases.



3.

(a) The pattern below (on page 10) should be the same. As we only have 1 point here. It does not change for increasing the flexibility.

(b) The codes begin on page 11, the graph is on page 12.

```
[1]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import make_pipeline

np.random.seed(1)  # Set seed for reproducibility

n = 1  # Number of observations

# Generate training data
x = np.sort(np.random.uniform(-3, 3, n))
# Linear Function:  $y = 2 + 3x + \text{noise}$ 
y = 2 + 3 * x + 2 * np.random.randn(n) #adding random noise, times 2 to make it
    ↪ more noisier

# Generate test data
x0 = np.sort(np.random.uniform(-3, 3, n))
y0 = 2 + 3 * x0 + 2 * np.random.randn(n)

# Define the degrees of polynomial features
degrees = range(1, 14)

# Initialize lists to store training and test errors
train_errors = []
test_errors = []

# Iterate over different degrees
for degree in degrees:
    # Fit model
    lm = make_pipeline(PolynomialFeatures(degree=degree), LinearRegression())
    lm.fit(x.reshape(-1, 1), y)

    # Predictions on training data
    yhat_train = lm.predict(x.reshape(-1, 1))
    train_err = np.mean((y - yhat_train)**2)

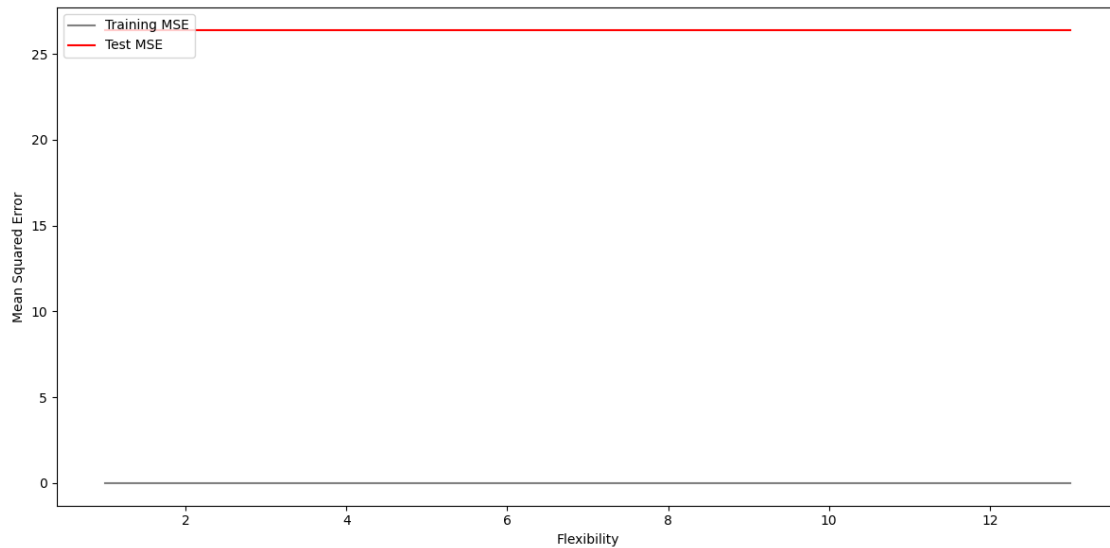
    # Predictions on test data
    yhat_test = lm.predict(x0.reshape(-1, 1))
    test_err = np.mean((y0 - yhat_test)**2)
```

```
# Append errors to lists
train_errors.append(train_err)
test_errors.append(test_err)

# Print errors
for degree, train_err, test_err in zip(degrees, train_errors, test_errors):
    print(f"Degree {degree}: Training error: {train_err}, Test error:␣
    ↪{test_err}")
```

```
Degree 1: Training error: 0.0, Test error: 26.396077576339028
Degree 2: Training error: 0.0, Test error: 26.396077576339028
Degree 3: Training error: 0.0, Test error: 26.396077576339028
Degree 4: Training error: 0.0, Test error: 26.396077576339028
Degree 5: Training error: 0.0, Test error: 26.396077576339028
Degree 6: Training error: 0.0, Test error: 26.396077576339028
Degree 7: Training error: 0.0, Test error: 26.396077576339028
Degree 8: Training error: 0.0, Test error: 26.396077576339028
Degree 9: Training error: 0.0, Test error: 26.396077576339028
Degree 10: Training error: 0.0, Test error: 26.396077576339028
Degree 11: Training error: 0.0, Test error: 26.396077576339028
Degree 12: Training error: 0.0, Test error: 26.396077576339028
Degree 13: Training error: 0.0, Test error: 26.396077576339028
```

```
[2]: # Plot
plt.figure(figsize=(12, 6))
plt.xlabel('Flexibility')
plt.ylabel('Mean Squared Error')
plt.plot(degrees, train_errors, color='grey', label= 'Training MSE')
plt.plot(degrees, test_errors, color='red', label= 'Test MSE')
plt.legend(loc='upper left')
plt.tight_layout()
plt.show()
```



```
[13]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import make_pipeline

np.random.seed(1) # Set seed for reproducibility

n = 1000 # Number of observations

# Generate training data
x = np.sort(np.random.uniform(-3, 3, n))
# Linear Function:  $y = 2 + 3x + \text{noise}$ 
y = 2 + 3 * x + 2 * np.random.randn(n) #adding random noise, times 2 to make it
    ↪ more noisier

# Generate test data
x0 = np.sort(np.random.uniform(-3, 3, n))
y0 = 2 + 3 * x0 + 2 * np.random.randn(n)

# Define the degrees of polynomial features
degrees = range(1, 50)

# Initialize lists to store training and test errors
train_errors = []
test_errors = []

# Iterate over different degrees
```

```

for degree in degrees:
    # Fit model
    lm = make_pipeline(PolynomialFeatures(degree=degree), LinearRegression())
    lm.fit(x.reshape(-1, 1), y)

    # Predictions on training data
    yhat_train = lm.predict(x.reshape(-1, 1))
    train_err = np.mean((y - yhat_train)**2)

    # Predictions on test data
    yhat_test = lm.predict(x0.reshape(-1, 1))
    test_err = np.mean((y0 - yhat_test)**2)

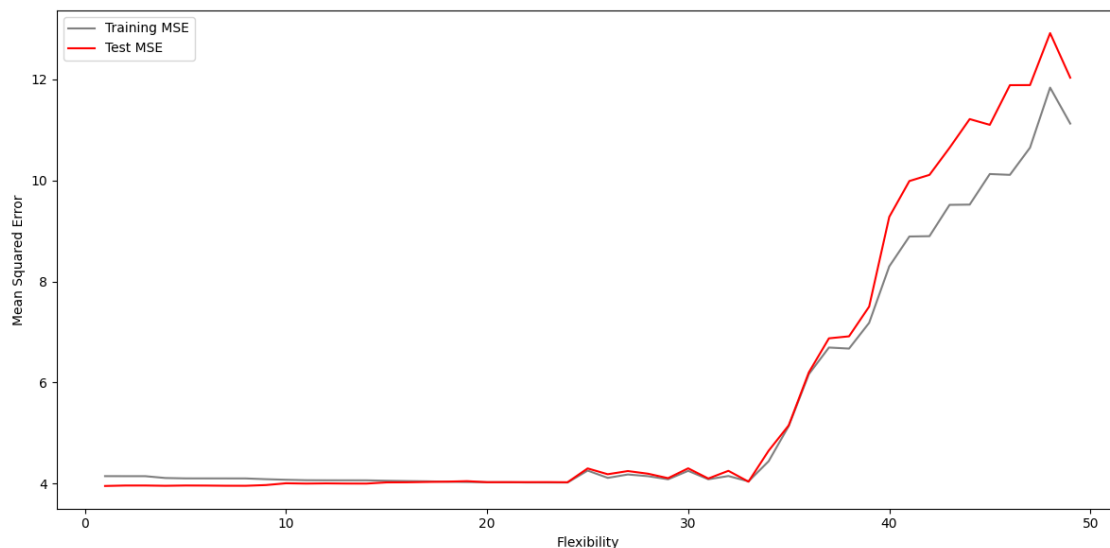
    # Append errors to lists
    train_errors.append(train_err)
    test_errors.append(test_err)

```

```

[14]: # Plot
plt.figure(figsize=(12, 6))
plt.xlabel('Flexibility')
plt.ylabel('Mean Squared Error')
plt.plot(degrees, train_errors, color='grey', label= 'Training MSE')
plt.plot(degrees, test_errors, color='red', label= 'Test MSE')
plt.legend(loc='upper left')
plt.tight_layout()
plt.show()

```



## Problem 2

1.

As  $p(x|\mu, \nu) := (\frac{\nu}{\mu})^\nu \frac{x^{\nu-1}}{\Gamma(\nu)} \exp\left(-\frac{\nu x}{\mu}\right)$ ,  $p(x|\theta) = (\frac{\nu}{\mu})^\nu \frac{x^{\nu-1}}{\Gamma(\nu)} \exp\left(-\frac{\nu x}{\mu}\right)$ .

$$L(\theta) = \prod_{i=1}^n p(x_i|\theta) = (\frac{\nu}{\mu})^{n\nu} \frac{\prod_{i=1}^n x_i^{\nu-1}}{\Gamma(\nu)^n} \exp\left(-\sum_{i=1}^n \frac{\nu x_i}{\mu}\right)$$

Suppose  $l(\theta)$  to be log-likelihood.

$$l(\theta) = \log L(\theta) = n\nu(\log \nu - \log \mu) + (\nu - 1) \sum_{i=1}^n \log x_i - n \log \Gamma(\nu) - \frac{\nu}{\mu} \sum_{i=1}^n x_i$$

2.

$$\frac{\partial l}{\partial \mu} = -\frac{1}{\mu} n\nu + \frac{\nu}{\mu^2} \sum_{i=1}^n x_i$$

$$\text{When } \frac{\partial l}{\partial \mu} = 0 \implies \frac{\sum x_i}{\mu} = n \implies \hat{\mu} = \frac{\sum x_i}{n} = \bar{x}$$

Checking the mean and variance and MLE, from what we've learned from Gaussian,  $\hat{\mu} = \bar{x}$  is not a surprise.

3.

$$\frac{\partial l}{\partial \nu} = n(\log \nu - \log \mu) + n - n \frac{\Gamma'(\nu)}{\Gamma(\nu)} + \sum_{i=1}^n \log x_i - \frac{\sum x_i}{\mu}$$

$$\text{When } \frac{\partial l}{\partial \nu} = 0 \implies n(\log \nu - \log \mu) + n - n \frac{\Gamma'(\nu)}{\Gamma(\nu)} + \sum_{i=1}^n \log x_i - n = n(\log \nu - \log \mu) +$$

$$\sum_{i=1}^n \log x_i - n \frac{\Gamma'(\nu)}{\Gamma(\nu)} = 0 \text{ where } n(\log \nu - \log \mu) + \sum_{i=1}^n \log x_i = \sum \log\left(\frac{x_i \hat{\nu}}{\mu}\right), -n \frac{\Gamma'(\nu)}{\Gamma(\nu)} = -\sum \phi(\hat{\nu}).$$

$$\text{So, } \sum_{i=1}^n (\log\left(\frac{x_i \hat{\nu}}{\mu}\right) - (\frac{x_i}{\mu} - 1) - \phi(\hat{\nu})) = n(\log \nu - \log \mu) + \sum_{i=1}^n \log x_i - \frac{\sum x_i}{\mu} + n - n \frac{\Gamma'(\nu)}{\Gamma(\nu)} = 0$$

## Problem 3

Our target is to get  $\min R(f)$  from the Bayes-Optimal Classifier  $f_0$ , due to the monotonicity of the integral  $R(f) = \int_{\mathbb{R}^d} R(f|\mathbf{x})p(\mathbf{x})d\mathbf{x}$ , we can say  $f_0$  minimizes  $R(f)$  by minimizing  $R(f|\mathbf{x}) := \sum_{y \in [K]} L^{0-1}(y, f(\mathbf{x}))P(y|\mathbf{x})$ . To find the  $f$  minimizes  $R(f|\mathbf{x})$ :

$$f(x) = \arg \min_{y \in [K]} \sum_{y \in [K]} L^{0-1}(y|f(\mathbf{x}))P(y|\mathbf{x})$$

$$= \arg \min_{y \in [K]} (1 - P(y|\mathbf{x}))$$

(mismatch loss is 1 and math loss 0

$$\implies 0 \cdot p(\text{match}) + 1 \cdot p(\text{mismatch}) = 1 \cdot (1 - p(\text{match})) = 1 - p(y|x)$$

So in this case,  $f(x) = \arg \max_{y \in [K]} P(y|\mathbf{x})$ , this is exactly  $f_0(\mathbf{x}) = \arg \max_{y \in [K]} P(y|\mathbf{x})$ . Hence  $f_0$  is what we want here.

## Problem 4

1.  $\text{posterior} \propto \text{likelihood} * \text{prior}$

As  $L(\theta_1, \dots, \theta_K | x_1, \dots, x_n) = \prod \theta^{n_k}$  and  $q(\theta_1, \dots, \theta_K) = \frac{1}{B(\alpha_1, \dots, \alpha_K)} \prod \theta^{\alpha_k - 1}$  where  $B(\alpha_1, \dots, \alpha_K) = \frac{\prod \Gamma(\alpha_k)}{\Gamma(\sum \alpha_k)}$ , so  $\Pi(\theta_1, \dots, \theta_K | x_1, \dots, x_n) = \prod \theta^{n_k} \cdot \frac{\Gamma(\sum \alpha_k)}{\prod \Gamma(\alpha_k)} \prod \theta^{\alpha_k - 1} = \frac{\Gamma(\sum_{k=1}^K \alpha_k)}{\prod_{k=1}^K \Gamma(\alpha_k)} \prod_{k=1}^K \theta^{\alpha_k + n_k - 1}$

2.

$$\begin{aligned}
\theta_{MAP}^{\hat{}} &= \arg \max \frac{\Gamma(\sum_{k=1}^K \alpha_k)}{\prod_{k=1}^K \Gamma(\alpha_k)} \prod_{k=1}^K \theta^{\alpha_k + n_k - 1} \\
&= \arg \max \prod_{k=1}^K \theta^{\alpha_k + n_k - 1} \\
&= \arg \max \log \prod_{k=1}^K \theta^{\alpha_k + n_k - 1} = \arg \max \sum_{k=1}^K \log \theta^{\alpha_k + n_k - 1} \\
&= \arg \max \sum_{k=1}^K (n_i + \alpha_i - 1) \log \theta_i \\
0 &= \frac{\partial}{\partial \theta_i} \left( \sum_{k=1}^K \log \theta^{\alpha_k + n_k - 1} + (n_k + \alpha_k - 1) \log \left( 1 - \sum_{i=1}^{K-1} \theta_i \right) \right) \\
&= \frac{(n_i + \alpha_i - 1)}{\theta_i} - \frac{(n_k + \alpha_k - 1)}{1 - \sum_{i=1}^{K-1} \theta_i} \\
&= \frac{(n_i + \alpha_i - 1)}{\theta_i} - \frac{(n_k + \alpha_k - 1)}{\theta_k}
\end{aligned}$$

That is  $\hat{\theta}_i = \frac{n_i + \alpha_i - 1}{\sum (n_i + \alpha_i - 1)}$